

Aim - Build a simple linear regression model to predict house prices based on features like the number of bedrooms and square footage

Import Statements and Their Purposes

1. `import numpy as np`

Purpose: Provides support for large, multi-dimensional arrays and matrices, along with a collection of mathematical functions to operate on these arrays.

2. `import pandas as pd`

Purpose: Used for data manipulation and analysis, including reading data from files like CSV and handling DataFrame operations.

3. `import matplotlib.pyplot as plt`

Purpose: Essential for creating visualizations such as line charts, scatter plots, and histograms.

4. `from sklearn.model_selection import train_test_split`

Purpose: Provides a method to split the dataset into training and testing subsets, ensuring a controlled and reproducible division.

5. `from sklearn.preprocessing import StandardScaler`

Purpose: Used for feature scaling, ensuring all numerical features have the same scale, typically necessary for machine learning algorithms sensitive to feature magnitudes.

6. `from sklearn.linear_model import LinearRegression`

Purpose: Contains the implementation of linear regression, a simple and widely-used predictive modeling algorithm.

7. `import seaborn as sns`

Purpose: Used for creating visually appealing statistical plots, such as scatter plots with regression lines or heatmaps.

Functions Used and Their Purposes

1. `pd.read_csv()`

- Reads a CSV file into a Pandas DataFrame.
- Allows for easy manipulation and analysis of structured data.

2. `dataset.iloc[:, :-1].values`

- Extracts all columns except the last one (independent variables) as a NumPy array.
- This isolates the features (X).

3. **dataset.iloc[:, -1].values**
 - Extracts the last column (dependent variable or target variable) as a NumPy array (y).
4. **train_test_split(X, y, test_size=0.2, random_state=42)**
 - Splits the dataset into training (80%) and testing (20%) sets.
 - The random_state ensures reproducibility of the split.
5. **StandardScaler()**
 - Initializes the standard scaler, which normalizes features by removing the mean and scaling to unit variance.
6. **sc.fit_transform(X_train)**
 - Computes the scaling parameters (mean and standard deviation) from the training set and applies scaling to it.
7. **sc.transform(X_test)**
 - Applies the scaling parameters (computed from training data) to the test data, ensuring consistent scaling.
8. **LinearRegression()**
 - Initializes the linear regression model.
9. **regressor.fit(X_train, y_train)**
 - Trains the linear regression model on the training data (X_train and y_train).
10. **regressor.predict(X_test)**
 - Uses the trained model to predict outcomes for the test data.
11. **plt.figure(figsize=(10, 6))**
 - Initializes a new figure for the plot with specified dimensions.
12. **sns.scatterplot()**
 - Creates a scatter plot to visualize the relationship between the actual and predicted prices.
13. **plt.plot()**
 - Plots a reference line representing perfect prediction ($y = x$).
14. **plt.title(), plt.xlabel(), plt.ylabel(), plt.grid()**
 - Adds a title, axis labels, and gridlines to the plot for better readability.

15. **sc.transform(test_input)**
 - Scales the new input data using the same parameters as the training data.
 16. **regressor.predict(scaled_input)**
 - Predicts the house price for the scaled input.
 17. **print()**
 - Outputs the predicted price in a human-readable format.
-

Why Each Step is Necessary

1. **Data Reading and Preparation:**
 - Essential for loading and isolating the features and target variables for analysis.
2. **Splitting the Dataset:**
 - Ensures the model is evaluated on unseen data, avoiding overfitting.
3. **Feature Scaling:**
 - Brings all features to the same scale, crucial for ensuring the algorithm performs optimally and weights features appropriately.
4. **Training the Model:**
 - Builds the predictive relationship between features (X_train) and the target variable (y_train).
5. **Prediction and Visualization:**
 - Provides insights into the model's performance and allows you to predict new, unseen values.
6. **Single Input Prediction:**
 - Demonstrates how to use the trained model to make predictions for custom data.